

Anuj Sadani

PRINCIPAL ENGINEER | AI/ML SYSTEMS | TECHNICAL LEADERSHIP

Pune, India | +91 9766880844 | anuj.k.sadani@gmail.com | anujsadani.in | linkedin.com/in/anujsadani | github.com/asadani

PROFILE

Principal engineer with 17 years of software development experience, including the last 3-4 years building and deploying applied ML and GenAI systems. Leads technical direction and people development while staying hands-on with production AI architecture, from model routing and fine-tuning to agent tooling and evaluation. Experience spans document intelligence, distributed systems, cloud platforms, and engineering standards.

CORE STRENGTHS

Technical leadership: architecture, technical roadmaps, hiring, mentoring, performance development, team structure, engineering standards

Platforms: Python, FastAPI, AWS, Bedrock, SageMaker, Docker, Kubernetes/EKS, Terraform, Databricks

Applied AI systems: ML/LLM routing, RAG, agentic systems, MCP, evaluation, fine-tuning, model serving

Practical tooling: LangChain, LangGraph, OpenRouter, LiteLLM

EXPERIENCE

Infrd

Bengaluru, India

Principal Software Development Engineer | May 2025-Present

ML Engineering Technical Architect | Nov 2021-May 2025

Technical Specialist | Jan 2021-Oct 2021

Technical direction and people development for 12+ engineers across three teams; engineering standards and AI/ML architecture for document extraction.

- Led technical direction and people development across three teams, including hiring 8-10 engineers, mentoring, supporting promotions, and helping restructure teams.
- Own the GenAI and ML architecture for a document-extraction platform processing about 1M pages per day, at roughly 1-4 cents per page depending on document type.
- Designed and led a hybrid ML/LLM architecture: smaller models handle routine extraction and LLMs handle harder cases, reducing processing cost by 30-40% compared with an LLM-only approach.
- Designed an MCP-based agent tooling layer for mortgage document processing, translating underwriting-domain rules into production tools with subject-matter experts.
- Fine-tuned small open models for entity extraction and integrated model-based processing into production workflows; set organization-wide Python practices for security, application patterns, and code quality.

Delivery Hero

Berlin, Germany

Software Engineer II | Feb 2020-Dec 2020

- Built and scaled Python/Flask services for real-time order processing in a distributed logistics platform.
- Worked with product and engineering leadership on objectives, roadmaps, and delivery plans; partnered with data science and BI teams to put analytical workflows into production.
- Contributed to reliability and incident management using Terraform, Helm, Grafana, and New Relic.

NVIDIA

Pune, India

Senior Tools Development Engineer | 2014-2020

Tools Development Engineer | 2010-2014

Software Quality Assurance Engineer | 2009-2010

- Led an 8+ person team building embedded automation frameworks and internal engineering tools supporting GPU and SoC development.
- Owned tools through requirements, architecture, implementation, deployment, and long-term maintenance.
- Improved engineering productivity through automation-first practices and introduced a crowdsourced contribution model for internal tooling.

SELECTED RESEARCH

Tool Attention Is All You Need - Dynamic tool gating and lazy schema loading for scalable agentic workflows. arXiv, 2026.

Locale-Conditioned Few-Shot Prompting for On-Device PII Substitution - Small-language-model PII substitution. arXiv, 2026.

The Great Recalibration - AI application architecture and industrialized AI services. SSRN, 2026.

PATENTS

Three pending US patent applications for document-entity extraction and query optimization: 19/375,298; 19/375,285; 19/432,186.

EDUCATION & CERTIFICATIONS

B.E., Information Technology
Government College of Engineering, Amravati | 2005-2009

Certified ScrumMaster, Scrum Alliance
(Credential ID 001022708)

IBM Enterprise Design Thinking, Co-Creator

Resume source: [codex-resume/resume.md](#)